

HW Assignment 6 (Due by 10:30am on Nov 16)

1 Implementation (150 points)

Implement the 3 versions of the perceptron algorithm discussed in class: the perceptron, the average perceptron, and the kernel perceptron. The algorithms should stop after achieving convergence, or after a predefined number of epochs T , whichever comes first. Make sure that you organize your code in folders as shown in the table below. Write code only in the Python files indicated in bold.

```
ml4900/  
  hw06/  
    code/  
      perceptron.py  
      exercise5.py  
      spam_exercise.py  
      newsgroups_exercise.py  
      spam_model_p.txt  
      spam_model_ap.txt  
      newsgroups_model_p1.txt  
      newsgroups_model_ap1.txt  
      newsgroups_model_p2.txt  
      newsgroups_model_ap2.txt  
      output.txt  
    data/  
      spam/  
        spam_train.txt  
        spam_train_svm.txt  
        spam_test.txt  
        spam_test_svm.txt  
        spam_vocab.txt  
      newsgroups/  
        stopwords.txt  
        newsgroups_train1.txt  
        newsgroups_test2.txt  
        newsgroups_train2.txt  
        newsgroups_test2.txt  
        newsgroups_vocab.txt
```

1. **Perceptron Convergence:** Validate experimentally the conclusions you reached for Exercise 5 in the previous assignment. For 5(a), show that the perceptron algorithm hits the same weight vector at different epochs during training and thus it will run forever. For 5(b), show the number of epochs the kernel perceptron needed to converge and the values of the dual parameters at the end. Make sure you process the training examples in the order given in the assignment. In **perceptron.py** you should implement these five functions:

- (a) `perceptron_train(data, labels, epochs)` trains the vanilla perceptron algorithm and returns the weight vector.
- (b) `aperceptron_train(data, labels, epochs)` trains the average perceptron algorithm and returns the average weight vector.
- (c) `perceptron_test(w, data)` test a perceptron with weights w and returns a vector with the labels on the test examples in data.
- (d) `kperceptron_train(data, labels, epochs, kernel)` trains the kernel perceptron algorithm and returns the vector α of parameters.
- (e) `kperceptron_test(alpha, data, kernel)` tests the kernel perceptron algorithm with parameters α and returns a vector with the labels on the test examples in data.

Write the code for validating Exercise 5 in the file **exercise5.py**.

2. **Spam vs. Non-Spam:** In this problem, you will train and evaluate spam classifiers using the perceptron and average perceptron algorithms. The dataset contains two files: `spam_train.txt` with 4,000 training examples and `spam_test.txt` with 1,000 test examples. The dataset is based on a subset of the SpamAssassin Public Corpus. Each line in the training and test files contains the pre-processed version of one email. The line starts with the label, followed by the email tokens separated by spaces.

Figure 1 shows a sample source email, while Figure 2 shows its pre-processed version in which web addresses are replaced with the “httpaddr” token, numbers are replaced with a “number” token, dollar amounts are replaced with “dollar numb”, and email addresses are replaced with “emailaddr”. Furthermore, all words are lower-cased, HTML tags are removed, and words are reduced to their stems i.e. “expecting”, “expected”, “expectation” are all replaced with “expect”. Non-words and punctuation symbols are removed.

> Anyone knows how much it costs to host a web portal ?
 > Well, it depends on how many visitors youre expecting. This can be anywhere from less than 10 bucks a month to a couple of \$100. You should checkout <http://www.rackspace.com/> or perhaps Amazon EC2 if youre running something big..

To unsubscribe yourself from this mailing list, send an email to: groupname-unsubscribe@egroups.com

Figure 1: Sample email from the SpamAssassin corpus.

anyon know how much it cost to host a web portal well it depend on how mani visitor your expect thi can be anywher from less than number buck a month to a coupl of dollar numb you should checkout httpaddr or perhap amazon ecnumb if your run someth big to unsubscrib yourself from thi mail list send an email to emailaddr

Figure 2: Pre-processed version of email from the SpamAssassin corpus.

- (a) Create a vocabulary file `spam_vocab.txt` that contains a list of only the (pre-processed) tokens that appear at least 30 times in the training examples. The file should contain one token per line in the format `<id> <token>`, where each token is associated a unique integer identifier. The tokens should be listed in increasing order of their identifiers, starting from 1. See for example the vocabulary file `newsgroups_vocab.txt` that we generated for the newsgroup classification problem. Implement the function that creates the vocabulary in `spam_exercise.py`.
 - (b) For each training and test example, create a sparse feature vector representation wherein each example is represented as one line in the file using the format `<label> <id1>:<val1> <id2>:<val2> ...`, where the id's are listed in increasing order and correspond only to tokens that appear in that example (use 1 for all values, representing that fact that the corresponding token appeared in the example). An example of this sparse representation can be seen in the file `newsgroups_train1.txt` that we generated for the newsgroup classification problem. Save the new version of the dataset in the files `spam_train_svm.txt` and `spam_test_svm.txt`. Implement the function that creates the sparse feature vector representations in `spam_exercise.py`.
 - (c) Write an additional function `read_examples(file_name)` inside `perceptron.py` that reads all examples from a file with sparse feature vectors and returns a tuple `(data, labels)` where the `data` is a two dimensional array containing all feature vectors, one per row, in the same order as in the file, and the `labels` is a vector containing the corresponding labels.
 - (d) Train the perceptron algorithm until convergence, by reading the training examples from `spam_train_svm.txt` and by calling `perceptron_train(data, labels, epochs)` in the code you write inside `spam_exercise.py`. Process the examples in the order they are listed in the files. Report the number of epochs needed for convergence and the total number of mistakes made and save the returned parameter vector in `spam_model_p.txt`. Test the perceptron algorithm by reading the parameter vector from `spam_model_p.txt` and the test examples from `spam_test_svm.txt` and calling `perceptron_test(w, data)` in the code inside `spam_exercise.py`. Report the test accuracy.
 - (e) Train the average perceptron algorithm until the corresponding perceptron convergences, by reading the training examples from `spam_train_svm.txt` and by calling `aperceptron_train(data, labels, epochs)` in the code inside `spam_exercise.py`. Process the examples in the order they are listed in the files. Report the number of epochs needed for convergence and the total number of mistakes made and save the returned average parameter vector in `spam_model_ap.txt`. Test the average perceptron algorithm by reading the parameter vector from `spam_model_ap.txt` and the test examples from `spam_test_svm.txt` and calling `perceptron_test(w, data)` in the code inside `spam_exercise.py`. Report the test accuracy. Compare the test accuracy between the perceptron and the average perceptron.
3. **Atheism vs. Religion:** In this problem, you will train and evaluate the binary perceptron and average perceptron algorithms on a subset of the 20 newsgroups dataset. In this subset, there are 857 positive example and 570 test examples, on the topics of atheism and religion. Newsgroup postings on the topic of Atheism (`alt.atheism`) are given

label 1, whereas newsgroup posting on the topic of Religion (`talk.religion.misc`) are given label 0. Thus, the models will be trained to distinguish between postings on Atheism and postings on Religion.

The feature vectors have already been created for you and are stored in files using the sparse feature vector representation described above. To create these feature vectors, we stripped metadata, quotes, and headers from the documents. The words were stemmed and tokens that appeared less than 20 times in the training examples were filtered out. Common tokens from the `stopwords.txt` file were also removed. The remaining tokens are stored in the vocabulary file `newsgroups_vocab.txt` and are used to create two versions of the dataset:

- **[Version 1]** In this version, each token corresponds to a feature whose value for a particular document is computed using the standard *tf.idf* formula from Information Retrieval (think search engines). The term frequency *tf* refers to the number of times the token appears in the document, whereas the inverse document frequency *idf* refers to the inverse of the log of the total number of documents that contain the token. The *idf* numbers are computed using the entire 20 newsgroup dataset. The two quantities are multiplied into one *tf.idf* value and are meant to give more importance to words that are rare (i.e. large *idf*) and appear more frequently inside the corresponding document example (i.e. large *tf*). The training and test examples thus created are stored in `newsgroups_train1.txt` and `newsgroups_test1.txt` respectively.
- **[Version 2]** This is the same as version 1 above, except that the term frequencies are set to 1 for all tokens that appear in a document, i.e. the number of times a token appears in the document is irrelevant and the only thing that matters is whether the token appeared or not in the document, and also how rare it is (through the *idf* weight). The training and test examples for this version created are stored in `newsgroups_train2.txt` and `newsgroups_test2.txt` respectively.

For more details on how the feature vectors were created, you can read the Scikit section at http://scikit-learn.org/stable/datasets/twenty_newsgroups.html.

For each version of the dataset, use the perceptron and average perceptron implementations that you wrote in `perceptron.py` and train the two algorithms for 10,000 epochs. Process the examples in the order they are listed in the files. Save the parameters in the corresponding `newsgroups_model-<x>.txt` files. Then evaluate the two perceptron algorithms on the corresponding test examples for each version. For each version, report and compare the accuracies of the two models.

4. **{Spam vs. Non-Spam} vs. {Atheism vs. Religion}**: Why is the accuracy on the second dataset (atheism vs. religion) much lower than the accuracy on the first dataset (spam vs. non-spam)? To answer this question accurately, you may want to download the original documents in the two datasets and look at some of them.

2 Submission

Turn in a hard copy of your homework report at the beginning of class on the due date. Electronically submit on Blackboard a hw06.zip file that contains the hw06 folder in which you write code **only in the required files**. The screen output produced when running the code should be redirected to (saved into) the **output.txt** file.

On a Linux system, creating the archive can be done using the command:

```
> zip -r hw06.zip hw06.
```

Please observe the following when handing in homework:

1. Structure, indent, and format your code well.
2. Use adequate comments, both block and in-line to document your code.
3. Make sure your code runs correctly when used in the directory structure shown above.